

## Alpaydin Chapter 2, Mitchell Chapter 7

- Alpaydin slides are in turquoise.
  - Ethem Alpaydin, copyright: The MIT Press, 2010.
  - `alpaydin@boun.edu.tr`
  - `http://www.cmpe.boun.edu.tr/~ethem/i2ml2e`
- All other slides are based on Mitchell.


# Learning a Class from Examples

- Class C of a “family car”
  - Prediction: Is car  $x$  a family car?
  - Knowledge extraction: What do people expect from a family car?
- Output:

Positive (+) and negative (–) examples
- Input representation:

$x_1$ : price,  $x_2$  : engine power

## Training Set $\mathcal{X}$




$$\mathcal{X} = \{\mathbf{x}^t, r^t\}_{t=1}^N$$


$$r = \begin{cases} 1 & \text{if } \mathbf{x} \text{ is positive} \\ 0 & \text{if } \mathbf{x} \text{ is negative} \end{cases}$$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$


## Class $\mathcal{C}$



## Hypothesis Class $\mathcal{H}$



## $S$ , $G$ , and the Version Space



most general hypothesis,  $G$

$h \in H$ , between  $S$  and  $G$  is  
consistent  
and make up the  
version space  
(Mitchell, 1997)

# Computational Learning Theory (from Mitchell Chapter 7)

- Theoretical characterization of the **difficulties** and **capabilities** of learning algorithms.
- Questions:
  - Conditions for successful/unsuccessful learning
  - Conditions of success for particular algorithms
- Two frameworks:
  - Probably Approximately Correct (PAC) framework: classes of hypotheses that can be learned; complexity of hypothesis space and bound on training set size.
  - Mistake bound framework: number of training errors made before correct hypothesis is determined.

# Computational Learning Theory

What general laws constrain inductive learning?

We seek theory to relate:

- Probability of successful learning
- Number of training examples
- Complexity of hypothesis space
- Accuracy to which target concept is approximated
- Manner in which training examples presented



## Specific Questions

- Sample complexity: How many training examples are needed for a learner to converge?
- Computational complexity: How much computational effort is needed for a learner to converge?
- Mistake bound: How many training examples will the learner misclassify before converging?


Issues: When to say it was successful? How are inputs acquired?

# Sample Complexity

How many training examples are sufficient to learn the target concept?

1. If learner proposes instances, as queries to teacher
  - Learner proposes instance  $x$ , teacher provides  $c(x)$
2. If teacher (who knows  $c$ ) provides training examples
  - teacher provides sequence of examples of form  $\langle x, c(x) \rangle$
3. If some random process (e.g., nature) proposes instances
  - instance  $x$  generated randomly, teacher provides  $c(x)$

# True Error of a Hypothesis



**Definition:** The **true error** (denoted  $error_{\mathcal{D}}(h)$ ) of hypothesis  $h$  with respect to target concept  $c$  and distribution  $\mathcal{D}$  is the probability that  $h$  will misclassify an instance drawn at random according to  $\mathcal{D}$ .

$$error_{\mathcal{D}}(h) \equiv \Pr_{x \in \mathcal{D}} [c(x) \neq h(x)]$$

## Two Notions of Error

*Training error* of hypothesis  $h$  with respect to target concept  $c$

- How often  $h(x) \neq c(x)$  over training instances

*True error* of hypothesis  $h$  with respect to  $c$


- How often  $h(x) \neq c(x)$  over future random instances

Our concern:

- Can we bound the true error of  $h$  given the training error of  $h$ ?
- First consider when training error of  $h$  is zero (i.e.,  $h \in VS_{H,D}$ )

# Exhausting the Version Space

Hypothesis space  $H$



( $r$  = training error,  $error$  = true error)

**Definition:** The version space  $VS_{H,D}$  is said to be  $\epsilon$ -**exhausted** with respect to  $c$  and  $\mathcal{D}$ , if every hypothesis  $h$  in  $VS_{H,D}$  has error less than  $\epsilon$  with respect to  $c$  and  $\mathcal{D}$ .

$$(\forall h \in VS_{H,D}) error_{\mathcal{D}}(h) < \epsilon$$

## How many examples will $\epsilon$ -exhaust the VS?

**Theorem:** [Haussler, 1988].

If the hypothesis space  $H$  is finite, and  $D$  is a sequence of  $m \geq 1$  independent random examples of some target concept  $c$ , then for any  $0 \leq \epsilon \leq 1$ , the probability that the version space with respect to  $H$  and  $D$  is not  $\epsilon$ -exhausted (with respect to  $c$ ) is less than

$$|H|e^{-\epsilon m}$$

This bounds the probability that any consistent learner will output a hypothesis  $h$  with  $error(h) \geq \epsilon$

If we want this probability to be below  $\delta$

$$|H|e^{-\epsilon m} \leq \delta$$

then

$$m \geq \frac{1}{\epsilon} (\ln |H| + \ln(1/\delta))$$

## Proof of $\epsilon$ -Exhausting Theorem

**Theorem:** Prob. of  $V S_{H,D}$  not being  $\epsilon$ -exhausted is  $\leq |H|e^{-\epsilon m}$ .

**Proof:**

- Let  $h_i \in H$  ( $i = 1..k$ ) be those that have true error greater than  $\epsilon$  wrt  $c$  ( $k \leq |H|$ ).
- We fail to  $\epsilon$ -exhaust the VS iff at least one  $h_i$  is consistent with all  $m$  sample training instances (note: they have true error greater than  $\epsilon$ ).
- Prob. of a single hypothesis with error  $> \epsilon$  is consistent for one random sample is at most  $(1 - \epsilon)$ .
- Prob. of that hypothesis being consistent with  $m$  samples is  $(1 - \epsilon)^m$ .
- Prob. of at least one of  $k$  hypotheses with error  $> \epsilon$  is consistent with  $m$  samples is  $k(1 - \epsilon)^m$ .
- Since  $k \leq |H|$ , and for  $0 \leq \epsilon \leq 1$ ,  $(1 - \epsilon) \leq e^{-\epsilon}$ :

$$k(1 - \epsilon)^m \leq |H|(1 - \epsilon)^m \leq |H|e^{-\epsilon m}$$

# PAC Learning

Consider a class  $C$  of possible target concepts defined over a set of instances  $X$  of length  $n$ , and a learner  $L$  using hypothesis space  $H$ .

*Definition:*  $C$  is **PAC-learnable** by  $L$  using  $H$  if for all  $c \in C$ , distributions  $\mathcal{D}$  over  $X$ ,  $\epsilon$  such that  $0 < \epsilon < 1/2$ , and  $\delta$  such that  $0 < \delta < 1/2$ ,

learner  $L$  will with probability at least  $(1 - \delta)$  output a hypothesis  $h \in H$  such that  $error_{\mathcal{D}}(h) \leq \epsilon$ , in time that is polynomial in  $1/\epsilon$ ,  $1/\delta$ ,  $n$  and  $size(c)$ .



# Agnostic Learning

So far, we assumed that  $c \in H$ . What if it is not the case?

Agnostic learning setting: don't assume  $c \in H$

- What do we want then?
  - The hypothesis  $h$  that makes fewest errors on training data
- What is sample complexity in this case?

$$m \geq \frac{1}{2\epsilon^2} (\ln |H| + \ln(1/\delta))$$

derived from Hoeffding bounds:

$$Pr[error_{\mathcal{D}}(h) > error_D(h) + \epsilon] \leq e^{-2m\epsilon^2}$$


# Shattering a Set of Instances

*Definition:* a **dichotomy** of a set  $S$  is a partition of  $S$  into two disjoint subsets.

*Definition:* a set of instances  $S$  is **shattered** by hypothesis space  $H$  if and only if for every dichotomy of  $S$  there exists some hypothesis in  $H$  consistent with this dichotomy.

# Three Instances Shattered

Instance space  $X$




Each closed contour indicates one dichotomy. What kind of hypothesis space  $H$  can shatter the instances?

# The Vapnik-Chervonenkis Dimension

*Definition:* The **Vapnik-Chervonenkis dimension**,  $VC(H)$ , of hypothesis space  $H$  defined over instance space  $X$  is the size of the largest finite subset of  $X$  shattered by  $H$ . If arbitrarily large finite sets of  $X$  can be shattered by  $H$ , then  $VC(H) \equiv \infty$ .

Note that  $|H|$  can be infinite, while  $VC(H)$  finite!

## VC Dim. of Linear Decision Surfaces



(a)



(b)

- When  $H$  is a set of lines, and  $S$  a set of points,  $VC(H) = 3$ .
- (a) can be shattered, but (b) cannot be. However, if at least one subset of size 3 can be shattered, that's fine.
- Set of size 4 cannot be shattered, for any combination of points (think about an XOR-like situation).

## VC Dimension: Another Example

$S = \{3.1, 5.7\}$ , and hypothesis space includes intervals  $a < x < b$ .

- Dichotomies: both, none, 3.1, or 5.7.
- Are there intervals that cover all the above dichotomies?

What about  $S = x_0, x_1, x_2$  for an arbitrary  $x_i$ ? (cf. collinear points).

## Sample Complexity from VC Dimension

How many randomly drawn examples suffice to  $\epsilon$ -exhaust  $VS_{H,D}$  with probability at least  $(1 - \delta)$ ?

$$m \geq \frac{1}{\epsilon} (4 \log_2(2/\delta) + 8VC(H) \log_2(13/\epsilon))$$

$VC(H)$  is directly related to the sample complexity:

- More expressive  $H$  needs more samples.
- More samples needed for  $H$  with more tunable parameters.

# Mistake Bounds

So far: how many examples needed to learn?

What about: how many mistakes before convergence?

- This is an interesting question because some learning systems may need to start operating while still learning.

Let's consider similar setting to PAC learning:

- Instances drawn at random from  $X$  according to distribution  $\mathcal{D}$ .
- Learner must classify each instance before receiving correct classification from teacher.
- Can we bound the number of mistakes learner makes before converging?



# Mistake Bounds: Halving Algorithm

Consider the Halving Algorithm:

- Learn concept using version space *Candidate-Elimination* or *List-Then-Eliminate* algorithm (no need to know details about these algorithms).
- Classify new instances by majority vote of version space members.

How many mistakes before converging to correct  $h$ ?

- ... in worst case?
- ... in best case?

## Mistake Bound of Halving Algorithm

- Start with version space =  $H$ .
- Mistake is made when more than half of the  $h \in H$  misclassified.
- In that case, at most half of  $h \in VS$  will be eliminated.
- That is, each **mistake** reduces the  $VS$  by half.
- Initially  $|VS| = |H|$ , and each mistake halves the  $VS$ , so it takes  $\log_2 |H|$  mistakes to reduce  $|VS|$  to 1.
- Actual worst-case bound is  $\lfloor \log_2 |H| \rfloor$ .

## Optimal Mistake Bounds

Let  $M_A(C)$  be the max number of mistakes made by algorithm  $A$  to learn concepts in  $C$ . (maximum over all possible  $c \in C$ , and all possible training sequences)

$$M_A(C) \equiv \max_{c \in C} M_A(c)$$

*Definition:* Let  $C$  be an arbitrary non-empty concept class. The **optimal mistake bound** for  $C$ , denoted  $Opt(C)$ , is the minimum over all possible learning algorithms  $A$  of  $M_A(C)$ .

$$Opt(C) \equiv \min_{A \in \text{learning algorithms}} M_A(C)$$

## Mistake Bounds and VC Dimension


Littlestone (1987) showed:

$$VC(C) \leq Opt(C) \leq M_{Halving}(C) \leq \log_2(|C|)$$


# Noise and Model Complexity

Use the simpler one because

- Simpler to use  
(lower computational complexity)
- Easier to train (lower space complexity)
- Easier to explain  
(more interpretable)
- Generalizes better (lower variance - Occam's razor)



## Multiple Classes, $C_i, i = 1, \dots, K$



$$\mathcal{X} = \{\mathbf{x}^t, r^t\}_{t=1}^N$$

$$r_i^t = \begin{cases} 1 & \text{if } \mathbf{x}^t \in C_i \\ 0 & \text{if } \mathbf{x}^t \in C_j, j \neq i \end{cases}$$

Train hypotheses  
 $h_i(\mathbf{x}), i = 1, \dots, K:$

$$h_i(\mathbf{x}^t) = \begin{cases} 1 & \text{if } \mathbf{x}^t \in C_i \\ 0 & \text{if } \mathbf{x}^t \in C_j, j \neq i \end{cases}$$

# Regression

- Learning is an ill-posed problem, data is not sufficient to find a unique solution
- The need for inductive bias, assumptions about  $\mathcal{H}$
- Generalization: How well a model performs on new data
- Overfitting:  $\mathcal{H}$  more complex than  $C$  or  $f$
- Underfitting:  $\mathcal{H}$  less complex than  $C$  or  $f$

# Model Selection & Generalization


$$\mathcal{X} = \{x^t, r^t\}_{t=1}^N$$

$$r^t \in \mathbb{R}$$

$$r^t = f(x^t) + \varepsilon$$

$$E(g | \mathcal{X}) = \frac{1}{N} \sum_{t=1}^N [r^t - g(x^t)]^2$$

$$E(w_1, w_0 | \mathcal{X}) = \frac{1}{N} \sum_{t=1}^N [r^t - (w_1 x^t + w_0)]^2$$





## Triple Trade-Off

- There is a trade-off between three factors (Dietterich, 2003):
  1. Complexity of  $\mathcal{H}$ ,  $c(\mathcal{H})$ ,
  2. Training set size,  $N$ ,
  3. Generalization error,  $E$ , on new data
- As  $N \uparrow$ ,  $E \downarrow$
- As  $c(\mathcal{H}) \uparrow$ , first  $E \downarrow$  and then  $E \uparrow$

## Cross-Validation

1. Model:  $g(\mathbf{x} | \theta)$

2. Loss function:  $E(\theta | \mathcal{X}) = \sum_t L(r^t, g(\mathbf{x}^t | \theta))$

3. Optimization procedure:

$$\theta^* = \operatorname{argmin}_{\theta} E(\theta | \mathcal{X})$$

# Dimensions of Supervised Learning

- To estimate generalization error, we need data unseen during training. We split the data as
  - Training set (50%)
  - Validation set (25%)
  - Test (publication) set (25%)
- Resampling when there is few data